

Latency-aware PM2.5 interval forecasting under simulated measurement and feedback delays: a reproducible Beijing benchmark

Renato Quispe Vargas^{1*}

¹Universidad Nacional del Altiplano.

Corresponding author(s). E-mail(s): rntvargasvav@gmail.com;

Abstract

Observation latency restricts both forecast inputs and responses available for recalibration. This retrospective Beijing benchmark evaluates 24-hour PM2.5 forecasts using 420,768 station-hour records and 102,465 observed final-period targets. Two distinct objectives organize the evidence. The original fixed-predictor analysis compares rolling calibration with valid feedback against static unmatched calibration under a fixed predictor. At a simulated 72-hour delay and 90% nominal level, rolling reduces mean interval score from 377.14 to 328.01, with coverage 88.32%; its paired difference interval excludes zero with seven-day blocks but includes zero with 14- and 28-day blocks. A post-review exploration evaluates alternative pipelines, including CQR with delay-specific training and an ACI-inspired daily delayed-feedback heuristic (ADH). At that same delay and nominal level, CQR has the lowest observed score among evaluated pipelines (258.40), coverage 88.91% and mean width 182.72, whereas ADH is closer to nominal coverage (89.78%) with greater width (203.17) and score (317.48). CQR coverage has a pointwise 14-day block-bootstrap interval of [86.51, 91.18]%; inclusion of 90% does not establish calibration. The CQR–rolling score reduction of 21.2% is exploratory and compares complete pipelines, not isolated calibration effects. In a residual-interval factorial, predictor adaptation reduces rolling score from 328.01 to 312.66. Randomized heterogeneous delays and distribution diagnostics assess sensitivity rather than independent confirmation. The contribution is a reproducible information-set benchmark, not a new algorithm, universal coverage guarantee or demonstrated transfer to other cities.

Keywords: PM2.5 forecasting, observation latency, uncertainty quantification, temporal validation, conformal calibration, environmental informatics

1 Introduction

Air-quality forecasts are made with the measurements available at their issue time. Retrospective datasets, in contrast, often provide a completed archive without an explicit record of publication latency. A forecast evaluation can therefore inadvertently grant access to measurements, or to forecast errors, before they would have become available. This distinction matters when prediction intervals are recalibrated online: the observed response needed to score an earlier forecast may itself be delayed.

Conformal inference offers a general route from predictive models to uncertainty sets. Adaptive conformal inference addresses distribution shift (Gibbs and Candès 2021), and subsequent work studies its behaviour and aggregation for dependent series (Zaffran et al. 2022). Sequential predictive conformal inference explicitly exploits dependence in conformity scores (Xu and Xie 2023). Missing covariates introduce another difficulty: average coverage can conceal poor coverage for particular missingness patterns (Zaffran et al. 2023). These established contributions motivate an empirical evaluation that keeps information availability explicit.

The present study uses the Beijing multi-site air-quality archive (Chen 2017; Zhang et al. 2017) to ask how point and interval forecasts change when the PM2.5 channel arrives late. The analysis does not propose a new conformal algorithm. Its contribution is a reproducible, temporally separated benchmark in which the same simulated delay governs both feature construction and eligibility of residuals for recalibration. An intentionally infeasible feedback rule, which ignores the latter restriction, is reported separately as a diagnostic.

Three questions guide the experiment. How much do simple point predictors deteriorate under increasing delay? Does matching the calibration information state or updating calibration over time improve interval performance? Are the resulting conclusions stable across locations, calendar periods and dependence-sensitive uncertainty calculations? Interval score is used alongside coverage and width to avoid rewarding intervals solely for being wide (Gneiting and Raftery 2007).

1.1 Related work and methodological positioning

Environmental forecasting already includes comparisons of conditional quantile models for NO₂ (Vasseur and Aznarte 2021) and multi-step prediction intervals for locally stationary pollutant series (Li et al. 2022). These studies are closer methodological precedents than generic artificial-intelligence reviews. The present question concerns a separate operational restriction: which inputs and forecast errors have arrived at issue time? Neither reference establishes the behaviour of the delayed-information procedures evaluated here.

The foundations of conformal prediction separate the construction of uncertainty sets from the choice of point learner (Vovk et al. 2005; Angelopoulos and Bates 2023). Regression methods include split-conformal inference and resampling-based alternatives such as the jackknife+ (Lei et al. 2018; Barber et al. 2021). Work beyond

exchangeability explicitly examines departures from the assumptions supporting classical guarantees (Barber et al. 2023); a recent survey organizes further developments from a data perspective (Zhou et al. 2026). These references motivate the benchmark but do not supply an automatic coverage theorem for its pooled, nonstationary environmental series. The original rolling residual rule is distinct from ACI, SPCI and jackknife+. The extension implements CQR and an ACI-inspired daily delayed-feedback heuristic (ADH), without attributing the original ACI theorem to that modification.

Evaluation design is a separate source of uncertainty. Forecasting research emphasizes out-of-sample assessment aligned with deployment (Tashman 2000). Cross-validation can be valid for particular autoregressive settings under appropriate conditions (Bergmeir et al. 2018); this does not justify random splitting for every dependent dataset. Structured environmental data require a splitting strategy that reflects the intended prediction task (Roberts et al. 2017). Selection on the same validation results used to estimate performance can induce optimism (Varma and Simon 2006), and information leakage undermines reproducibility in machine-learning-based science (Kapoor and Narayanan 2023). Accordingly, this study separates development, calibration and test periods and makes feedback arrival explicit.

2 Data and chronological design

2.1 Source and audit

The UCI dataset contains hourly air-pollutant and meteorological records at 12 Beijing sites from 1 March 2013 through 28 February 2017. It is distributed under CC BY 4.0 (Chen 2017). The analysis uses the twelve original station files inside the archived PRSA2017 data package. Auxiliary files named `data.csv` and `test.csv` in the outer download are not included. A source checksum and file manifest are retained in the reproducibility materials.

There are 420,768 station-hour records and 18 original columns. The audit found no duplicate station-time keys and no discontinuities in the hourly row grid. PM2.5 is missing in 8,739 records (2.08%). These absences form 2,837 contiguous station-specific episodes; the longest lasts 343 hours at Aotizhongxin, from 17 December 2014 at 10:00 through 31 December 2014 at 16:00. Across the archive, 81 hourly timestamps have no PM2.5 measurement at any of the twelve sites. These are descriptions of recorded availability, not confirmed sensor failures.

Missing-data mechanisms affect the interpretation of analyses on observed records (Rubin 1976). No missing-at-random mechanism is established for this archive. Raw concentrations are preserved without outlier deletion or response imputation. No negative values occur in the six pollutant channels, but this check does not establish instrumental accuracy. Predictor absences remain explicit. Time stamps retain the source clock; no time-zone conversion or real publication-time reconstruction is claimed. Audits of missingness used the complete archive before the experiment, whereas model selection and concentration summaries were restricted to development data.

2.2 Forecast target and sample flow

For site s and issue time t , the response is the hourly concentration $Y_{s,t+24}$, not a daily average. Model development covers March 2013–February 2015, interval calibration covers March 2015–February 2016, and the final test covers March 2016–February 2017. The last 24 issue times per site in each period are excluded so that their response does not cross the period boundary. The first 168 hours of the archive are reserved for constructing the longest historical window. A missing response prevents scoring that origin; a missing input does not automatically remove it.

Table 1 Sample flow. Exclusions are mutually exclusive and applied from left to right. Counts refer to station-specific forecast origins.

Period	Initial	Boundary	History	Response NA	Eligible
Development	210,240	288	2,016	4,226	203,710
Calibration	105,408	288	0	2,120	103,000
Test	105,120	288	0	2,367	102,465

The test includes origins from 1 March 2016 through 27 February 2017; their responses extend through 28 February. Observations from an earlier period can supply historical inputs for a later period, as they would in chronological deployment. No random train-test split is used. The same 102,465 test origins are evaluated under every delay, enabling paired comparisons. This conditions evaluation on originally observed responses and does not establish performance when the unobserved target values are systematically different.

3 Methods

3.1 Information sets and simulated delay

Let $d \in \{0, 6, 24, 72, 168\}$ denote a constant PM2.5 delivery delay in hours. At issue time t , the PM2.5 history is restricted to times $u \leq t - d$. Other pollutant and meteorological channels are assumed to be available at t , retaining their natural missingness. Thus the experiment represents a delayed pollutant channel, not loss of an entire station or a measured latency distribution. The zero-delay condition still contains all naturally missing values.

The baseline feature set contains station identity, current pollutants and meteorology, missingness indicators, PM2.5 lags of 1, 3, 6, 12, 24, 48 and 168 hours, retrospective means, standard deviations and observed counts over 24 and 168 hours, age of the most recent available PM2.5 value, and sinusoidal calendar encodings. Station and wind-direction categories are one-hot encoded using development categories. For $d > 0$, current PM2.5 is hidden, lags shorter than d are hidden, and all dependent

summaries are rebuilt. Specifically, a w -hour summary uses only

$$\{Y_{s,u} : t - w + 1 \leq u \leq t - d, Y_{s,u} \text{ observed}\}. \quad (1)$$

If this set is empty, its mean and standard deviation are missing and its count is zero. Age is computed from the latest observed time no later than $t - d$. Univariate summaries omit missing values and carry their effective count. They never interpolate from future measurements.

3.2 Point predictors and development selection

Ridge regularization and gradient boosting provide established statistical baselines (Hoerl and Kennard 1970; Friedman 2001). The cited boosting framework is conceptual background, not a claim that the original implementation is identical to the histogram estimator used here. In the original experiment, four model families are considered: last-observation persistence; a training mean by station and response hour; Ridge regression; and histogram gradient boosting with absolute-error loss. Persistence falls back to a station-specific median of training responses if no historical value is available. Ridge uses training-median imputation and standardization. Boosting handles numerical missing values directly. All point predictions are truncated below at zero.

Three expanding-origin validation windows cover March–June 2014, July–October 2014 and November 2014–February 2015. Each training set contains only responses dated before the corresponding window. The validation windows also exclude targets beyond their right boundary. Ridge penalties are 1, 10 and 100. Boosting configurations comprise 15 leaves and 100 iterations, 31 leaves and 100 iterations, and 31 leaves and 200 iterations, with learning rate 0.05, minimum leaf size 40 and L_2 regularization 1. Random early stopping is disabled. Configuration selection minimizes the unweighted mean of the three validation MAEs, separately within each family.

The selected models are refitted on the full development period, with no retraining during calibration or test. A fixed boosting model is used for the original interval comparisons, as specified before evaluation. The later exploratory extension is described separately. All four families remain in the point-forecast comparison even when rankings change under delay. The protocol is a locally recorded analysis plan, not a public preregistration.

3.3 Residual calibration with arrival restrictions

For a fixed predictor $\hat{f}_d$ evaluated with delayed features, let

$$R_{s,t,d} = |Y_{s,t+24} - \hat{f}_d(X_{s,t}^{(d)})|. \quad (2)$$

The notation does not imply separate predictor training by delay: the same fitted model receives the differently observed features. Given n available residuals and nominal level $1 - \alpha$, define q_α as the residual of rank $\lceil (n + 1)(1 - \alpha) \rceil$. The prediction

interval is

$$C_{s,t,d} = \left[\max\{0, \hat{Y}_{s,t+24,d} - q_\alpha\}, \hat{Y}_{s,t+24,d} + q_\alpha \right]. \quad (3)$$

Nominal levels are 90% and 95%. Residuals are pooled across stations. This finite-sample quantile convention follows conformal calibration, but dependence and distribution changes preclude asserting an automatic distribution-free guarantee here.

Static unmatched calibration uses residuals computed without an imposed delay, applying the same quantile to every delayed scenario. Static matched calibration recomputes the residuals under the scenario’s feature delay. Both use calibration responses before 23 February 2016 at 00:00, a common cutoff seven days before the first test origin that makes them available under all delays.

Rolling valid calibration is updated once per day. At the start of day a , it uses residuals with response measurement times v satisfying

$$a - d - 30 \text{ days} < v \leq a - d. \quad (4)$$

The associated forecast origins are at most $a - d - 24$ hours. Eligible residuals can come from the calibration period or earlier test forecasts once their responses arrive. Quantiles remain fixed within a day, while point forecasts change hourly. As an implementation diagnostic, a fourth method sets the feedback delay to zero while retaining the delayed input features. For $d > 0$ this rule is operationally infeasible and is excluded from claims about deployable methods.

3.4 Evaluation and uncertainty

Point metrics are MAE and RMSE. Their different loss scales can produce different rankings, making the evaluation criterion part of the forecasting question (Hyndman and Koehler 2006). Interval evaluation reports empirical coverage, mean width, and the interval score

$$\begin{aligned} \text{IS}_\alpha(l, u; y) &= (u - l) + \frac{2}{\alpha}(l - y)\mathbf{1}(y < l) \\ &\quad + \frac{2}{\alpha}(y - u)\mathbf{1}(y > u), \end{aligned} \quad (5)$$

which combines interval width and penalties for observations outside its limits (Gneiting and Raftery 2007). Smaller scores are better. Scores and widths have the same concentration units as the response. Marginal coverage and subgroup coverage are both reported; neither is inferred from interval score alone. This distinction follows the separation between calibration and sharpness in probabilistic forecast assessment (Gneiting et al. 2007). Threshold- and quantile-weighted scores offer additional ways to emphasize particular predictive regions (Gneiting and Ranjan 2011); those alternatives are not computed in this benchmark.

The original primary comparison is rolling valid minus static unmatched interval score at 90% nominal coverage and $d = 72$ hours. Daily sums of paired score differences and daily observation counts are resampled together, retaining all stations within each sampled day. Resampling contiguous blocks is motivated by the dependent-data

bootstrap literature (Kunsch 1989; Lahiri 2003). The stationary bootstrap instead uses random block lengths (Politis and Romano 1994); that alternative is not the procedure implemented here. A circular block bootstrap uses 1,000 replications, seven-day blocks, and percentile limits of 2.5% and 97.5%. Fourteen- and 28-day blocks assess sensitivity. The original generator is NumPy 2.3.5 `default_rng(20261003)`, using PCG64. A single generator advances continuously through block lengths 7, 14 and 28, in that order. Each replicate draws $\lceil N/L \rceil$ start days uniformly with replacement from $0, \dots, N-1$, expands blocks modulo N , and truncates to N days. Percentiles use `numpy.quantile` with its `linear` interpolation method on unrounded replicates; only displayed limits are rounded to two decimal places. These intervals condition on the fitted models and this network; they omit model-refitting uncertainty. Nonstationarity and the artificial year-end join of a circular bootstrap remain limitations. No hourly independence-based significance test is used.

Circular blocks were chosen in the local protocol to avoid preferentially excluding boundary days from block starts while preserving short sequences and all within-day station dependence. This is a convenient dependent-data approximation, not a model of seasonal stationarity. Its artificial February–March join and mixing of seasonal regimes can distort persistence and percentile limits in either direction. A non-circular moving-block bootstrap removes that join but changes boundary inclusion frequencies; a stationary bootstrap randomizes block lengths and does not itself remove seasonal nonstationarity. Resampling whole years is not informative with one final test year; four unequal, seasonally distinct quarter groups are too few for reliable percentile inference. Leaving a period out addresses stability under a changed evaluation sample rather than the same full-period sampling distribution. These alternatives were not implemented, and block-length sensitivity cannot validate the circular assumption.

3.5 Reproducibility and verification

The accompanying materials include source and protocol hashes, exact partitions, feature specifications, model selection results, fitted estimators, predictions for every delay and tabulated outcomes. Feature checks compare sampled lag values and retrospective means against direct time-key lookups. A separate R implementation recomputes coverage, width and interval score from saved predictions for all 40 original configurations: five delays $\times$ two nominal levels $\times$ four rules (static unmatched, static matched, rolling valid, and the rolling rule that ignores feedback delay), agreeing with Python within 10^{-8} . The primary paired mean is independently reconciled. A separate R check recomputed all 68 extension rows: 60 homogeneous-delay configurations (five delays $\times$ two levels $\times$ six methods, including CQR and ADH) and eight heterogeneous-delay configurations (two levels $\times$ four rules). These checks recompute metrics from saved bounds. Separately, all ten CQR interval configurations were reconstructed from the twenty saved quantile learners. ADH uses actual issued-interval errors with delayed arrivals; its daily feedback recursion is explicitly checked in the additional implementation audit described below. Python 3.12.14 and R 4.5.2 were used; package versions are recorded with the executable scripts. This verifies arithmetic consistency, not scientific validity of all assumptions.

Local protocol chronology and its limits.

Local session records place the protocol write request at 3 October 2026, 01:12:34.531 UTC; the file creation/write timestamp is 01:12:39.1505342 UTC. The selection file, carrying the protocol hash, was written at 01:15:19.6009483 UTC, before the first recorded final-test command at 01:15:49.776 UTC. The original interval table and final-run manifest were written at 01:17:08.2118407 and 01:17:08.5461841 UTC. The protocol bytes match the initial write payload. Its SHA-256 is:

78c518b2081dc77d4e987595b3fff99aea27958f8e40bb326da86b0e61131af4

Hash comparisons also show that the original interval table, primary bootstrap and daily paired differences match the public v0.2.0 and v0.3.0 snapshots. A supplementary provenance report records these checks. Local logs and filesystem dates are mutable: they document the observed workflow, not an independent timestamp certification or proof that no unrecorded analysis occurred. The entire observed sequence took approximately five minutes. No separate trusted freeze timestamp exists. We call this a local protocol recorded before execution, not preregistration or evidence of independent confirmation.

4 Original fixed-predictor results

P: Original analysis with a local protocol

Available local records place the original protocol before the first recorded test execution and identify one primary contrast. This is not public preregistration or independent external confirmation. Additional configurations are descriptive; sensitivity analyses are labelled S.

Table 2 Main analysis map: objectives, data and evidential status. P denotes the original analysis with a local protocol recorded before execution; it is not independent confirmation.

ID	Status	Data	Main comparison	Outcome
P	Local protocol	Final Beijing test	Rolling vs static unmatched, fixed predictor	Score reduction; sensitive to block length
E	Post-review exploratory	Same final test	CQR with delay-specific ADH	Different score/coverage/width trade-offs
S	Sensitivity	Same archive; earlier Beijing period	Block size, pool-ing, missingness and delay scenarios	Robustness and scope; no external replication

4.1 Temporal selection and point forecasts

The smallest boosting configuration achieves the lowest mean validation MAE (53.00), compared with 60.04 for persistence. Increasing tree complexity does not improve

the aggregate selection criterion (Table 3). Ridge with penalty 100 has the lowest validation RMSE, 70.79, although its MAE is higher than boosting. The selection therefore already displays a trade-off between error criteria.

Table 3 Mean development-validation errors across the three temporal windows, in $\mu\text{g}/\text{m}^3$. Configuration selection uses MAE within each family.

Model	MAE	RMSE
Boosting: 15 leaves, 100 iterations	53.00	73.67
Boosting: 31 leaves, 100 iterations	53.28	73.85
Boosting: 31 leaves, 200 iterations	53.72	74.26
Ridge: alpha 100	54.35	70.79
Ridge: alpha 10	54.36	70.80
Ridge: alpha 1	54.36	70.80
Station-hour mean	59.18	75.59
Persistence	60.04	84.45

On the final test without additional delay, boosting obtains MAE 51.27, versus 58.12 for persistence and 53.06 for Ridge. Ridge again has smaller RMSE, 73.91 versus 79.53 for boosting. With a 72-hour delay, boosting MAE increases to 55.66 and Ridge MAE is 53.44. Persistence deteriorates to 78.52. These comparisons concern a forecast 24 hours ahead even when its PM2.5 history is several days old. The other pollutant channels remain current by design, and may retain information correlated with the missing channel.




Fig. 1 Point-forecast errors under the five PM2.5 delivery delays. Models are trained once on development data and evaluated on identical test origins. The ordering of models depends on the error criterion and delay.

4.2 Interval performance under delay

At $d = 72$ hours and 90% nominal coverage, the static unmatched interval has coverage 88.15%, width 169.20 and score 377.14. Static matching changes these to 88.41%, 170.64 and 375.39. Rolling valid calibration attains coverage 88.32%, width 184.09 and score 328.01 (Table 4). Its average score is 13.03% lower than static unmatched calibration, but its intervals are wider and its coverage remains below 90%. Thus neither matching nor rolling updating establishes nominal coverage.

Table 4 P: Original fixed-predictor comparison at a 72-hour delay. Coverage is expressed as a percentage; width and score are in $\mu\text{g}/\text{m}^3$. Each row uses 102,465 targets.

Nominal	Method	Coverage	Width	Score
90	Static unmatched	88.15	169.20	377.14
90	Static matched	88.41	170.64	375.39
90	Rolling valid	88.32	184.09	328.01
95	Static unmatched	93.59	221.15	458.87
95	Static matched	94.58	236.13	438.42
95	Rolling valid	93.14	225.21	391.54

At 95% nominal coverage, static matching achieves 94.58% coverage and score 438.42, whereas rolling valid calibration attains 93.14% and score 391.54. Better average score therefore coexists with a larger coverage deficit. Across the five delays, rolling calibration improves the average score relative to the static unmatched reference, but the experiment does not demonstrate a universal ordering of calibration quality (Fig. 2).




Fig. 2 Coverage, width and interval score at the 90% nominal level. Delay values are discrete experimental scenarios shown at equally spaced positions. The dashed line marks nominal coverage; all three approaches fall below it in aggregate.

4.3 Sensitivity within the original analysis (S)

The infeasible rolling diagnostic gives score 324.82 rather than 328.01 at 90% and a 72-hour delay. At 95% with a 168-hour delay it gives 385.01 rather than 399.95, making performance appear about 3.74% better if response arrival is ignored. These are descriptive differences from one test year. They do not justify assuming that premature feedback improves every metric or every subgroup.

For rolling valid intervals at 90% and $d = 72$, station-level coverage ranges from 85.60% at Dongsi to 91.93% at Huairou. Calendar-quarter coverage ranges from 81.91% in 2016Q4 to 95.29% in 2017Q1. The first and last calendar quarters are partial: 2016Q1 contains March only, and 2017Q1 contains January and February test origins. Subgroup results are descriptive and were not used to retune the evaluated models.




Fig. 3 Coverage heterogeneity for arrival-aware rolling intervals at a 72-hour delay and a 90% nominal level. The dashed lines indicate nominal coverage. Partial first and last quarters and unequal numbers of observed targets limit direct seasonal comparisons.

The primary paired score difference is -49.13. Its seven-day block-bootstrap interval excludes zero narrowly, but the intervals using 14- and 28-day blocks include zero (Table 5; Fig. 4). Consequently the result is suggestive of a lower average score in this sample, not robust evidence of superiority across reasonable block lengths. Treating 102,465 dependent test records as independent would not reflect this sensitivity.

Table 5 S: Sensitivity of the primary paired score comparison to bootstrap block length. Negative values favour rolling valid calibration over static unmatched calibration.

Block (days)	Difference	95% percentile interval
7	-49.13	[-107.13, -3.02]
14	-49.13	[-118.54, 6.24]
28	-49.13	[-133.34, 3.60]




Fig. 4 Paired interval-score difference and block-bootstrap limits at 90% nominal coverage and a 72-hour delay. All estimates use the same observed targets and fixed fitted predictor.

5 Post-review exploratory extension

E: Exploratory comparisons

All added comparator analyses and their paired bootstrap intervals were specified after inspection of the principal test results. They cannot be promoted to independent confirmatory evidence.

5.1 Analysis status and training regimes

The following analyses were designed after the original test results had been inspected. Their separate local protocol and outputs preserve this chronology; they are exploratory additions, not a new preregistered test. No configuration is selected by final-test score. Experiment A retains the original fixed predictor. Experiment B selects 15 or 31 leaves, with 100 boosting iterations, using the same three expanding development folds but reconstructing predictors under the candidate delay. Training responses must have arrived before the relevant fold boundary. The final adapted fit uses only development responses arriving before 1 March 2015. Columns entirely unobserved in a training fold are excluded using that fold only. MAE summarizes absolute

Table 6 Information and selection regimes. The idealized reference is not deployable when measurements are delayed.

Regime	Training information	Selection and interpretation
A: fixed	Original, zero imposed delay	Original development choice; sensitivity to information loss
B: adapted	Same delay as evaluation	Delay-specific temporal validation; arrival-purged responses
Idealized	Zero imposed delay	Reference on the same targets; not a delayed deployment option

loss, whereas RMSE gives greater weight to large errors. The original boosting model is fitted with absolute-error loss; its MAE advantage does not establish superiority for decisions dominated by extremes. Station-level point metrics are supplied alongside Table 7.

5.2 Exact executed CQR algorithm

For each $d \in \{0, 6, 24, 72, 168\}$ hours and nominal error $\alpha \in \{0.10, 0.05\}$, let $X_t^{(d)}$ denote the causal delayed predictor vector. Four separate histogram-boosting quantile models $f_{d,p}$, $p \in \{.025, .05, .95, .975\}$, use pinball loss $\rho_p(u) = u\{p - \mathbb{I}(u < 0)\}$ (Koenker and Bassett 1978; Romano et al. 2019). Training uses development responses satisfying $t + 24 + d < 1$ March 2015, with training-only feature exclusion for entirely unobserved columns. Each quantile model has 15 leaves, 100 iterations, learning rate .05, minimum leaf size 40, L_2 regularization 1, no early stopping and seed 20261003.

Table 7 Fixed and delay-adapted point predictors. MAE and RMSE are in $\mu\text{g}/\text{m}^3$.

Delay	Fixed MAE	Adapted MAE	Fixed RMSE	Adapted RMSE
0	51.27	51.27	79.53	79.53
6	54.76	51.47	82.71	79.80
24	55.26	51.67	84.39	79.81
72	55.66	51.29	83.39	78.44
168	55.51	51.04	84.01	78.12

These settings are fixed: the 15/31-leaf candidate selection described in Section 5.1 applies to the point predictor, not to these quantile models.

First clip each predicted quantile at zero, $g_{d,p}(x) = \max\{0, f_{d,p}(x)\}$, then order the pair:

$$\ell_d(x) = \min\{g_{d,\alpha/2}(x), g_{d,1-\alpha/2}(x)\}, \quad (6)$$

$$u_d(x) = \max\{g_{d,\alpha/2}(x), g_{d,1-\alpha/2}(x)\}. \quad (7)$$

The calibration set $\mathcal{C}$ comprises naturally observed calibration targets with measurement time strictly before 23 February 2016. The cutoff $c_\star = 23$ February 2016, 00:00 precedes the first test issue $t_\star = 1$ March 2016, 00:00 by exactly 168 hours (2016 is a leap year). Writing v_i for target measurement time, every retained response satisfies

$$v_i < c_\star, \quad 0 \leq d \leq 168 \text{ h} \implies r_i = v_i + d < c_\star + 168 \text{ h} = t_\star. \quad (8)$$

The latest retained hourly target is 22 February at 23:00, arriving at 29 February at 23:00 for $d = 168$ h, one hour before $t_\star$. Of the 103,000 naturally observed calibration responses surviving the original split rules, 1,988 are excluded by this common target-time cutoff. Therefore $n = |\mathcal{C}| = 101,012$ counts observed station-specific responses, with identical keys at every delay. No additional row is removed for missing predictors: the delayed reconstruction preserves rows, missing numeric values are handled by boosting, and categorical encodings are fixed from training. All 45 required source columns exist under every delay; entirely unobserved training columns can be excluded from the encoded design without deleting responses. The fitted designs retain 73, 70, 66, 64 and 62 encoded columns for delays 0, 6, 24, 72 and 168 h, respectively. These column exclusions do not change the response count. The supplementary sample/feature audit lists the retained encoded columns and a common hash of the calibration keys. There are twenty quantile learners (four per delay), combined into ten interval configurations (two nominal levels per delay); all use the same calibration-response exclusions. Compute the *signed* scores and rank correction:

$$s_i = \max\{\ell_d(X_i) - Y_i, Y_i - u_d(X_i)\}, \quad k = \lceil (n+1)(1-\alpha) \rceil, \quad Q_{d,\alpha} = s_{(k)}. \quad (9)$$

Here $s_{(k)}$ is the k th smallest score, without interpolation. Negative scores are retained; scores are neither absolute-valued nor clipped. The executed implementation asserts $1 \leq k \leq n$ and $Q_{d,\alpha} \geq 0$. The latter holds for every fitted configuration; inputs violating these assertions require an explicit methodological amendment rather than silent substitution. The issued CQR interval is

$$C_{t,d,\alpha} = \left[\max\{0, \ell_d(X_t^{(d)}) - Q_{d,\alpha}\}, u_d(X_t^{(d)}) + Q_{d,\alpha} \right]. \quad (10)$$

The nonnegative lower truncation uses the nonnegative concentration support. Both models and correction remain fixed throughout testing: test responses never recalibrate CQR. The raw quantile baseline uses $[\ell_d, u_d]$. At 72 hours, the corrections are 1.93848 and 1.25839 for 90% and 95%; their ranks are 90,912 and 95,963. All ten configurations and all stored test endpoints were reconstructed from the saved quantile models, with maximum discrepancy below 6×10^{-14} . Crossing counts were zero in this execution, although ordering remains part of the defined algorithm. The audit table supplies all corrections and arrival cutoffs. Serial dependence and distribution shift preclude importing an exchangeable-sample coverage guarantee. Contrasts with the fixed residual rules compare complete training-and-calibration pipelines, not calibration alone.

5.3 ACI-inspired heuristic: name and scope

We call the adaptive comparator the *ACI-inspired daily delayed-feedback heuristic (ADH)* (Gibbs and Candès 2021; Zaffran et al. 2022). It is not a reproduction of unmodified ACI and carries no transferred ACI guarantee. The historical machine identifier `ACI_delayed_projected` is retained only for compatibility with the published prediction files.

Let a_j be a daily midnight, r_i the response arrival time, and C_i^{issued} the interval actually issued for an earlier test forecast at t_i . Define

$$F_j = \{i : a_{j-1} < r_i \leq a_j, t_i < a_j, i \text{ is a test forecast}\}, \quad (11)$$


$$E_j = |F_j|^{-1} \sum_{i \in F_j} \mathbb{1}\{Y_i \notin C_i^{\text{issued}}\}, \quad (12)$$

$$\eta_j = \Pi_{[.001, .5]} \{\eta_{j-1} + .01(\alpha - E_j)\}. \quad (13)$$

Initialize $\eta_0 = \alpha$; skip the update when F_j is empty, including the first test update. Interval boundaries count as covered. A day receives one update regardless of its arriving batch size. Next take the rank-corrected $(1 - \eta_j)$ quantile of absolute residuals with arrivals in $(a_j - 30 \text{ days}, a_j]$, pooled across stations, and issue nonnegative-truncated symmetric intervals around the fixed point predictor for that day. Arrivals at the exact midnight are processed before issuance. Daily batching, unequal batch sizes, projection, a rolling residual window and delayed feedback differ from the original ACI setting. Neither gain .01 nor bounds $[.001, .5]$ were optimized on the test. These choices define an exploratory heuristic; they establish neither long-run nor conditional coverage. SPCI, EnbPI and missing-data augmentation remain unimplemented.

5.4 Information-set audit

For constant delay d , a forecast issued at t concerns Y_{t+24} , whose arrival is $t + 24 + d$. The first permitted daily update is the next midnight at or after arrival, assuming arrivals at that exact instant are processed before issuance. At each update, the residual set contains only arrivals in $(a - 30 \text{ days}, a]$. The fitted predictors remain fixed; only the calibration quantile and, for ADH, the error-level state evolve. Event ledgers



At update a : use only residuals with arrival $\leq a$; keep fitted predictors fixed.

Fig. 5 Measurement, arrival and update are distinct events. The same delay restricts predictors and feedback.

retain issue time, target time, arrival, first inclusion, last possible use and whether the residual is used within the evaluated period. Daily audit tables record the latest eligible arrival. Independent checks reconstruct selected quantiles and perturb unavailable residuals to verify invariance. For heterogeneous delays, eligibility uses each station's own arrival time rather than a shared measurement cutoff.

Table 8 E: Exploratory pipeline trade-offs at 72 hours. CQR denotes CQR with delay-specific training. Coverage and its pointwise 95% circular 14-day block-bootstrap interval are percentages (1,000 replicates; no simultaneous multiplicity adjustment). Width and score are in $\mu\text{g}/\text{m}^3$.

Nominal	Method	Coverage	Coverage interval	Width	Score
90	CQR	88.91	[86.51, 91.18]	182.72	258.40
90	ADH	89.78	[87.15, 92.42]	203.17	317.48
90	Rolling, by station	88.32	[85.30, 91.12]	184.94	324.86
90	Rolling, equal weights	88.32	[85.23, 91.16]	184.08	328.05
95	CQR	93.82	[92.03, 95.44]	211.01	307.18
95	ADH	94.64	[92.68, 96.37]	253.79	367.12
95	Rolling, by station	93.28	[91.10, 95.30]	225.69	386.58
95	Rolling, equal weights	93.13	[90.85, 95.21]	225.16	391.70

At 90%, CQR's observed coverage is below nominal. Its dependence-aware interval [86.51, 91.18]% contains 90%, so the point estimate alone does not establish a distinguishable departure from nominal; neither does interval inclusion prove calibration.

CQR has the lowest observed score among the evaluated pipelines at 72 hours, at both nominal levels. ADH is closer to nominal coverage but produces wider intervals and a larger score. These are distinct criteria: coverage measures inclusion, width describes concentration, and the chosen score penalizes both width and misses. No uniform or application-independent superiority follows.

5.5 Paired score uncertainty for exploratory contrasts

The exploratory paired contrasts use 10,000 replicates to reduce Monte Carlo error in percentile limits; the original protocol retains its originally specified 1,000-replicate analysis. The previously reported exploratory coverage and subgroup intervals also retain 1,000 replicates and seed 20261005 + L . A larger replicate count improves numerical precision of the bootstrap approximation, not its assumptions or evidential status. For the direct comparisons, 10,000 circular day-block draws resample paired score sums and observed-target counts together, retaining all stations per day. Blocks are 7, 14 or 28 days; the random seed is 20261006 + L for block length L . Each bootstrap mean is the resampled total difference divided by the resampled total target count, preserving the observation-weighted estimand. Pointwise percentile limits are the .025 and .975 quantiles. These intervals condition on stored fitted models; they include neither model-selection/refitting uncertainty nor simultaneous multiplicity control. Circular joining and nonstationarity remain limitations. This separate extension does not replace the original primary bootstrap.

Table 9 E: Exploratory paired score differences at 72 hours. Pointwise 95% block-bootstrap intervals; no simultaneous multiplicity adjustment. Negative values favour the first method. Units: $\mu\text{g}/\text{m}^3$.

Nominal	Contrast	Difference	7-day interval	14-day interval	28-day interval
90	CQR – rolling	-69.61	[-112.27, -32.38]	[-113.16, -30.66]	[-119.36, -24.91]
90	ADH – rolling	-10.53	[-31.74, 4.70]	[-30.98, 5.24]	[-33.94, 6.71]
90	CQR – ADH	-59.07	[-99.96, -21.60]	[-99.84, -21.36]	[-101.59, -22.69]
95	CQR – rolling	-84.36	[-157.06, -21.09]	[-156.06, -21.72]	[-155.84, -22.91]
95	ADH – rolling	-24.42	[-59.11, 3.36]	[-56.75, 2.90]	[-63.42, 5.41]
95	CQR – ADH	-59.95	[-125.99, 6.27]	[-125.36, 1.90]	[-120.84, -10.60]

This table compares entire pipelines: CQR changes the fitted quantile learner, delay-specific training, interval shape and calibration relative to the original fixed-point-predictor rolling rule. It does not isolate the causal effect of conformal calibration or predictor adaptation. At 90%, the descriptive score reduction is $69.61/328.01 \approx 21.2\%$. Here rolling means the original global rolling-valid rule (score 328.01), not the equal-station-weight sensitivity (328.05).

At 90%, CQR minus rolling has mean difference -69.61 and 14-day interval $[-113.16, -30.66]$. ADH minus rolling has difference -10.53 with interval $[-30.98, 5.24]$; an adaptive-heuristic advantage is therefore not established by this

comparison. CQR minus ADH is -59.07 , with interval $[-99.84, -21.36]$. Even intervals excluding zero remain post-hoc, exploratory evidence from the same inspected test period. The supplementary contrast table includes all delays, both nominal levels, static baselines, raw quantile intervals, and separately labelled earlier-period and heterogeneous-delay sensitivities.

5.6 Exploratory predictor-by-calibrator factorial

To separate predictor adaptation from recalibration, an additional post-review analysis crosses two stored point predictors (original fixed and delay-adapted) with three residual rules (static delay-matched, rolling global, and ADH). The full design has $5 \times 2 \times 2 \times 3 = 60$ configurations across delays and nominal levels. No model is refitted or selected on test performance. Every cell uses the same observed test responses and nonnegative-truncated symmetric residual interval. Static calibration uses the same 101,012 response keys before 23 February 2016, with residuals recomputed from the relevant predictor. Rolling and ADH use the same arrival eligibility, 30-day window and daily schedule; ADH retains its step size and projection. Thus within-predictor contrasts change only the calibration rule; within-calibrator contrasts change the predictor and its corresponding residuals. This controlled algorithm comparison is conditional on these fitted models, not a randomized causal experiment.

Table 10 E: Predictor-by-calibrator factorial at 72 h. Coverage is a percentage; width and score are in $\mu\text{g}/\text{m}^3$. All cells use 102,465 identical observed test targets. These are point estimates; paired score uncertainty is reported below. Static means delay-matched calibration, not the original static-unmatched primary comparator.

Nominal	Predictor	Calibrator	Coverage	Width	Score
90	Fixed	Static matched	88.41	170.64	375.39
90	Fixed	Rolling	88.32	184.09	328.01
90	Fixed	ADH	89.78	203.17	317.48
90	Adapted	Static matched	88.32	166.48	346.75
90	Adapted	Rolling	88.03	175.63	312.66
90	Adapted	ADH	89.64	192.68	303.20
95	Fixed	Static matched	94.58	236.13	438.42
95	Fixed	Rolling	93.14	225.21	391.54
95	Fixed	ADH	94.64	253.79	367.12
95	Adapted	Static matched	93.76	211.44	424.09
95	Adapted	Rolling	93.12	213.79	370.04
95	Adapted	ADH	94.67	239.05	350.18

At 90% nominal level, changing from the fixed to the adapted predictor changes score by -28.65 $[-48.88, -10.31]$ under static calibration, -15.35 $[-24.67, -6.55]$ under rolling, and -14.27 $[-22.22, -6.86]$ under ADH. Brackets here are pointwise 95% intervals

using 14-day blocks. Holding the predictor fixed, rolling minus static matched is -47.38 [-115.73, 3.69]; with the adapted predictor it is -34.09 [-87.54, 6.12].

Define the interaction for rule $r \in \{\text{rolling, ADH}\}$ as

$$I_r = (S_{\text{adapted},r} - S_{\text{adapted,static}}) - (S_{\text{fixed},r} - S_{\text{fixed,static}}), \quad (14)$$

where S is mean interval score. Negative interaction means a larger score reduction from that rule under the adapted predictor. For rolling, I_r is 13.29 [-1.25, 30.61]; for ADH it is 14.37 [-2.70, 34.20] at 90%. Both interaction intervals include zero, so the analysis does not establish an additional interaction advantage. These contrasts show why a single CQR-versus-rolling difference cannot be allocated entirely to calibration. CQR remains a comparison of whole pipelines: this factorial does not separately identify its quantile loss, learner, asymmetric interval shape and conformal correction.

The supplement reports 270 paired score contrasts: nine contrasts per delay/level and three block lengths. Each uses 10,000 circular day-block replicates, lengths 7/14/28, PCG64 seed 20261007 + L , observation-weighted paired daily sums, and linear-interpolated percentile limits. Inference is pointwise, conditional on fitted models and without multiplicity adjustment. Model refitting and independent test replication are not represented. R separately recomputed all 60 metric rows and exact covered counts within 10^{-8} . All 30 fixed-predictor cells reproduced the original saved bounds to numerical precision. At zero imposed delay the stored fixed and adapted predictors coincide, providing a zero-effect consistency check.

6 Sensitivity analyses and scope diagnostics

S: Sensitivity evidence

These checks characterize dependence assumptions, pooling, missingness and scope. The original block-length sensitivity is preserved; post-review sensitivities use already inspected data. The earlier Beijing period is not independent external validation.

6.1 Coverage uncertainty and station pooling

The supplementary tables report coverage, paired coverage differences and paired score differences using 1,000 circular day-block bootstrap replicates at block lengths 7, 14 and 28. Each resampled day retains all stations. For station and quarter comparisons at 72 hours, exploratory simultaneous bands use the bootstrap maximum absolute deviation of coverage differences within each method-level subgroup family. These do not establish global multiplicity control across every reported analysis. Replicates without observations in every subgroup are excluded and their count is reported. Conditional-on-fit inference and nonstationarity remain limitations.

At 72 h and 90%, global, station-specific and equal-station rolling each cover exactly 90,493 of 102,465 targets (88.3160103450%). This equality is not caused by rounding and does not imply identical intervals or covered cases. Relative to global pooling, station-specific pooling gains 1,048 covered cases and loses 1,048; equal-station weighting gains three and loses three. At 95%, their coverages differ

(93.1420485044%, 93.2757526960% and 93.1303371883%). The supplementary audit records exact counts and disagreement sets.

For the original rolling rule at 90%, the 14-day coverage interval is [85.23, 91.17]%; its coverage difference from static unmatched is 0.17 percentage points, with interval [-3.54, 3.94]. Figure 6 shows that coverage estimates should not be interpreted without dependence-aware uncertainty. Inclusion of nominal coverage in an approximate interval is not proof of calibration.




Fig. 6 Coverage uncertainty at 72 hours and 90% nominal level. Intervals are pointwise 95% circular 14-day block-bootstrap intervals, conditional on fitted models.

Pooling is examined using a common global residual quantile, separate station quantiles, and a weighted empirical global quantile giving each station equal total weight within the active window. The weighted rule is a sensitivity analysis, not a claim of exact finite-sample conformal validity. More observations give a station greater influence under ordinary pooling, whereas equal total weights remove that particular imbalance. All three methods retain temporal and cross-station dependence.

6.2 Practical loss, missingness and scope

In the original 72-hour, 90% comparison, rolling improves the daily mean score on 56.6% of days and produces wider intervals for 47.3% of targets. It newly covers 4,217 targets but loses coverage of 4,043, a net change of 174. The mean width change is $14.88 \mu\text{g}/\text{m}^3$, and the change in the score's exclusion penalty is -64.01. Thus the score improvement cannot be equated with a uniform gain in inclusion. Daily score-difference quantiles, concentration-stratified comparisons and hypothetical width-plus-exclusion-distance losses are provided. These losses are not monetary costs or validated operational utilities.

Missingness is quantified by station, calendar month, hour, study period and development-defined quintiles of PM10, NO2, temperature, pressure and wind speed. All absent PM2.5 records are retained in this audit. Concentration-stratified performance uses development-defined response quantiles and is retrospective; the target

concentration is unknown at issuance. The audit does not identify the missingness mechanism. Observation-probability weighting is not reported as a correction because it would require an additional observation model, positivity assessment and unverifiable assumptions about absent outcomes. Population-wide performance remains unidentified without such assumptions.

One additional synthetic scenario assigns delays 6, 24, 72 and 168 hours cyclically to alphabetically ordered stations. This tests heterogeneous availability across the network; it does not reproduce intermittent, out-of-order or concentration-dependent delivery. An earlier temporal backtest trains on responses arriving before January 2014, calibrates on later development records before 22 February 2015, and evaluates 103,000 targets in March 2015–February 2016 at a 72-hour delay. At 90%, static matched, rolling, ADH and CQR have scores 306.23, 296.79, 295.23 and 263.41, respectively; their coverages are 92.69%, 89.66%, 89.79% and 92.94%. The model configuration was fixed for this sensitivity. Its data overlap the broader study’s development/calibration history, so this is neither external replication nor independent confirmation.

6.3 Randomized delay assignments and distribution diagnostics

The alphabetical heterogeneous assignment ties station identity to delay and is therefore only one constructed scenario. We repeated the analysis with 40 random balanced assignments (PCG64 seed 20261007), each allocating three stations to each of 6, 24, 72 and 168 h. The same stored fixed predictor is used throughout, with features and residual arrival times corresponding to each assigned delay. All four pooling/adaptive rules and both nominal levels are evaluated in every assignment. Assignment means and dispersion describe sensitivity to this synthetic design; they are not confidence intervals or independent temporal replications. An optimized arrival-sorted implementation was checked against the earlier daily implementation for the alphabetical assignment and the first random assignment, across all eight method/level combinations, including the ADH feedback recursion. At 90%, global rolling has assignment-mean score 329.63 (SD 0.26; range 329.12–330.19), versus 329.63 for the alphabetical assignment. ADH has mean score 318.64 (SD 0.34; range 317.83–319.35). Mean coverages are 88.19% and 89.97%, respectively. Assignment-specific metrics and mappings are supplied so the average does not hide station-delay variation.

Distribution diagnostics compare naturally observed calibration and test responses and the absolute errors of the fixed and adapted predictors at 72 h (Table 11; Fig. 7). We report empirical quantiles, standardized mean differences, two-sample empirical CDF distances (KS), and Wasserstein distances, overall and for matched calendar months. No iid-based KS significance test is used. Quarter summaries retain the partial March and January–February endpoint quarters. These are descriptive comparisons of temporally dependent, selectively observed data; they do not prove a particular shift mechanism or explain all coverage deficits.

Table 11 S: Calibration-to-test distribution diagnostics at 72 h. Means and 90th percentiles are in $\mu\text{g}/\text{m}^3$; KS is the maximum absolute empirical CDF difference, without an iid-based p-value.

Variable	Cal. mean	Test mean	Cal. $q_{.90}$	Test $q_{.90}$	KS
Response	73.19	77.98	169.00	185.00	0.040
Fixed absolute error	54.68	55.66	101.46	112.64	0.018
Adapted absolute error	47.91	51.29	100.84	112.35	0.032




Fig. 7 S: Empirical distributions for calibration and test responses and 72-hour absolute prediction errors. Curves use naturally observed targets and fixed stored models. Horizontal axes are truncated at the pooled 99th percentile for readability; computations use all values. AI-assisted plotting code operates on archived numerical outputs.

For the fixed predictor, the 90th percentile of absolute error changes from 101.46 to 112.64. A frozen residual threshold therefore faces a different test-period error distribution. The monthly and quarterly diagnostics assess the seasonal variation hidden by the annual average. These differences are compatible with temporal distribution change, but missing-target selection, imperfect residual intervals and pooling also limit interpretation; annual undercoverage alone cannot establish distribution shift.

7 Discussion

Three observations stand out. First, temporal validation does not support a simple claim that more complex boosting configurations are preferable. Ridge and boosting trade places depending on the metric and delay. This argues for retaining transparent baselines and several error criteria when reporting an environmental forecasting model.

Second, the central difficulty is not only the loss of a recent predictor. The available residual sample also changes when the target channel arrives late. The diagnostic results quantify an apparent improvement that can arise when this constraint is ignored, especially at long delays and higher nominal coverage. A reproducible benchmark should therefore specify an arrival time for every measurement entering its update rule, not merely avoid future covariates in the original feature matrix.

Third, calibration choices involve trade-offs. Rolling residuals produce better average scores here, but their 95% coverage is lower than the delay-matched static method.

Static matching partly compensates for predictor degradation without guaranteeing nominal coverage. Strong variation among quarters and stations further cautions against treating the network-wide average as a property shared by every location or future period. These findings are consistent with the broader distinction between marginal and more local notions of coverage in missing-data inference (Zaffran et al. 2023).

The original objective tests whether recalibrating a fixed predictor with valid delayed feedback changes score relative to static unmatched calibration. The exploratory objective asks how alternative complete pipelines change that ranking. At 72 hours, CQR with delay-specific training has the lowest observed score among evaluated pipelines; ADH is closer to nominal coverage at the cost of greater width. These findings do not identify a universally preferable method. A lower score for rolling relative to a static residual rule does not establish it as the best evaluated pipeline. The factorial in Section 5.6 estimates recalibration contrasts separately within fixed and adapted point predictors, and predictor contrasts within each residual rule. Conditional quantile estimation and interval shape remain additional differences for CQR. The experiment still does not compare the full range of time-series or missing-data methods (Xu and Xie 2023; Zaffran et al. 2023).

Several design choices constrain interpretation. Simulated constant delays do not reproduce intermittent network failures, outcome-dependent missingness, or co-occurring loss of other channels. Reusing current correlated pollutants can soften the effect of delayed PM2.5. Natural missing responses are not recovered, so scoring is conditional on observation. There is only one city, one principal final test year and an overlapping earlier temporal sensitivity. Delay-adapted training improves the design but does not exhaust the attainable model space. All post-review comparisons are exploratory because the principal test was already examined.

The bootstrap intervals condition on the fitted predictor and assume that resampling time blocks gives a useful approximation despite seasonal changes. They omit uncertainty from model selection and refitting, and cannot turn twelve correlated stations into independent regional replications. Sensitivity to block length is therefore a substantive result rather than a technical detail. The evidence is insufficient to claim general improvement, operational readiness, or a new distribution-free guarantee.

8 Conclusion

The original fixed-predictor comparison shows a lower average interval score for rolling-valid than static-unmatched calibration under the original fixed predictor, with 88.32% coverage at 90% nominal and a superiority conclusion sensitive to bootstrap block length. This result addresses feedback-aware recalibration, not the optimal forecasting pipeline.

The post-review exploration shows that delay-specific training and conditional quantile estimation can substantially change the observed ranking. At 72 hours, CQR with delay-specific training has the lowest observed score among the evaluated pipelines, while ADH is closer to nominal coverage with wider intervals. CQR still undercovers in its point estimate, and approximate intervals do not establish

calibration. These are whole-pipeline, exploratory contrasts on an already inspected test set. The contribution is an auditable benchmark that separates information availability, score, coverage and width. External replication, independent confirmation and application-specific loss assessment remain necessary before deployment recommendations.

Supplementary information

The local research package contains the data-preparation script and original archive, chronological split indices, the locally recorded protocol, development scores, all test predictions, subgroup and bootstrap tables, and the independent R verification script. The research materials are provided through the GitHub repository listed below; the frozen analysis-code DOI is listed below.

References

- Angelopoulos, A.N., Bates, S.: Conformal prediction: A gentle introduction. *Foundations and Trends® in Machine Learning* **16**(4), 494–591 (2023) DOI: [10.1561/2200000101](https://doi.org/10.1561/2200000101)
- Barber, R.F., Candès, E.J., Ramdas, A., Tibshirani, R.J.: Predictive inference with the jackknife+. *The Annals of Statistics* **49**(1) (2021) DOI: [10.1214/20-aos1965](https://doi.org/10.1214/20-aos1965)
- Barber, R.F., Candès, E.J., Ramdas, A., Tibshirani, R.J.: Conformal prediction beyond exchangeability. *The Annals of Statistics* **51**(2) (2023) DOI: [10.1214/23-aos2276](https://doi.org/10.1214/23-aos2276)
- Bergmeir, C., Hyndman, R.J., Koo, B.: A note on the validity of cross-validation for evaluating autoregressive time series prediction. *Computational Statistics & Data Analysis* **120**, 70–83 (2018) DOI: [10.1016/j.csda.2017.11.003](https://doi.org/10.1016/j.csda.2017.11.003)
- Chen, S.: Beijing Multi-Site Air Quality. UCI Machine Learning Repository. Dataset; CC BY 4.0 (2017). DOI: [10.24432/c5rk5g](https://doi.org/10.24432/c5rk5g)
- Friedman, J.H.: Greedy function approximation: A gradient boosting machine. *The Annals of Statistics* **29**(5) (2001) DOI: [10.1214/aos/1013203451](https://doi.org/10.1214/aos/1013203451)
- Gneiting, T., Balabdaoui, F., Raftery, A.E.: Probabilistic forecasts, calibration and sharpness. *Journal of the Royal Statistical Society Series B: Statistical Methodology* **69**(2), 243–268 (2007) DOI: [10.1111/j.1467-9868.2007.00587.x](https://doi.org/10.1111/j.1467-9868.2007.00587.x)
- Gibbs, I., Candès, E.: Adaptive conformal inference under distribution shift. In: *Advances in Neural Information Processing Systems*, vol. 34 (2021). DOI: [10.48550/arxiv.2106.00170](https://doi.org/10.48550/arxiv.2106.00170) . DOI links to the corresponding arXiv preprint, not the proceedings version. [Online source](#)

- Gneiting, T., Raftery, A.E.: Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association* **102**(477), 359–378 (2007) DOI: [10.1198/016214506000001437](https://doi.org/10.1198/016214506000001437)
- Gneiting, T., Ranjan, R.: Comparing density forecasts using threshold- and quantile-weighted scoring rules. *Journal of Business & Economic Statistics* **29**(3), 411–422 (2011) DOI: [10.1198/jbes.2010.08110](https://doi.org/10.1198/jbes.2010.08110)
- Hoerl, A.E., Kennard, R.W.: Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics* **12**(1), 55–67 (1970) DOI: [10.1080/00401706.1970.10488634](https://doi.org/10.1080/00401706.1970.10488634)
- Hyndman, R.J., Koehler, A.B.: Another look at measures of forecast accuracy. *International Journal of Forecasting* **22**(4), 679–688 (2006) DOI: [10.1016/j.ijforecast.2006.03.001](https://doi.org/10.1016/j.ijforecast.2006.03.001)
- Koenker, R., Bassett, G.: Regression quantiles. *Econometrica* **46**(1), 33 (1978) DOI: [10.2307/1913643](https://doi.org/10.2307/1913643)
- Kapoor, S., Narayanan, A.: Leakage and the reproducibility crisis in machine-learning-based science. *Patterns* **4**(9), 100804 (2023) DOI: [10.1016/j.patter.2023.100804](https://doi.org/10.1016/j.patter.2023.100804)
- Kunsch, H.R.: The jackknife and the bootstrap for general stationary observations. *The Annals of Statistics* **17**(3) (1989) DOI: [10.1214/aos/1176347265](https://doi.org/10.1214/aos/1176347265)
- Lahiri, S.N.: *Resampling Methods for Dependent Data*. Springer, New York, NY (2003). DOI: [10.1007/978-1-4757-3803-2](https://doi.org/10.1007/978-1-4757-3803-2)
- Lei, J., G’Sell, M., Rinaldo, A., Tibshirani, R.J., Wasserman, L.: Distribution-free predictive inference for regression. *Journal of the American Statistical Association* **113**(523), 1094–1111 (2018) DOI: [10.1080/01621459.2017.1307116](https://doi.org/10.1080/01621459.2017.1307116)
- Li, J., Hu, Q., Zhang, F.: Multi-step-ahead prediction interval for locally stationary time series with application to air pollutant concentration data. *Stat* **11**(1), 411 (2022) DOI: [10.1002/sta4.411](https://doi.org/10.1002/sta4.411)
- Politis, D.N., Romano, J.P.: The stationary bootstrap. *Journal of the American Statistical Association* **89**(428), 1303–1313 (1994) DOI: [10.1080/01621459.1994.10476870](https://doi.org/10.1080/01621459.1994.10476870)
- Roberts, D.R., Bahn, V., Ciuti, S., Boyce, M.S., Elith, J., Guillera-Arroita, G., Hauenstein, S., Lahoz-Monfort, J.J., Schröder, B., Thuiller, W., Warton, D.I., Wintle, B.A., Hartig, F., Dormann, C.F.: Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography* **40**(8), 913–929 (2017) DOI: [10.1111/ecog.02881](https://doi.org/10.1111/ecog.02881)
- Romano, Y., Patterson, E., Candès, E.: Conformalized quantile regression. In: *Advances in Neural Information Processing Systems*, vol. 32 (2019). DOI: [10.48550/arxiv.1905.03222](https://doi.org/10.48550/arxiv.1905.03222) . DOI links to the corresponding arXiv preprint, not

- the proceedings version. [Online source](#)
- Rubin, D.B.: Inference and missing data. *Biometrika* **63**(3), 581–592 (1976)
DOI: [10.1093/biomet/63.3.581](https://doi.org/10.1093/biomet/63.3.581)
- Tashman, L.J.: Out-of-sample tests of forecasting accuracy: an analysis and review. *International Journal of Forecasting* **16**(4), 437–450 (2000)
DOI: [10.1016/s0169-2070\(00\)00065-0](https://doi.org/10.1016/s0169-2070(00)00065-0)
- Vasseur, S.P., Aznarte, J.L.: Comparing quantile regression methods for probabilistic forecasting of NO₂ pollution levels. *Scientific Reports* **11**, 11592 (2021)
DOI: [10.1038/s41598-021-90063-3](https://doi.org/10.1038/s41598-021-90063-3)
- Vovk, V., Gammerman, A., Shafer, G.: *Algorithmic Learning in a Random World*. Springer, New York, NY (2005). DOI: [10.1007/b106715](https://doi.org/10.1007/b106715)
- Varma, S., Simon, R.: Bias in error estimation when using cross-validation for model selection. *BMC Bioinformatics* **7**(1), 91 (2006) DOI: [10.1186/1471-2105-7-91](https://doi.org/10.1186/1471-2105-7-91)
- Xu, C., Xie, Y.: Sequential predictive conformal inference for time series. In: *Proceedings of the 40th International Conference on Machine Learning. Proceedings of Machine Learning Research*, vol. 202, pp. 38707–38727 (2023). DOI: [10.48550/arxiv.2212.03463](https://doi.org/10.48550/arxiv.2212.03463) . DOI links to the corresponding arXiv preprint, not the proceedings version. [Online source](#)
- Zhou, X., Chen, B., Gui, Y., Cheng, L.: Conformal prediction: A data perspective. *ACM Computing Surveys* **58**(2), 1–37 (2026) DOI: [10.1145/3736575](https://doi.org/10.1145/3736575)
- Zaffran, M., Dieuleveut, A., Josse, J., Romano, Y.: Conformal prediction with missing values. In: *Proceedings of the 40th International Conference on Machine Learning. Proceedings of Machine Learning Research*, vol. 202, pp. 40578–40604 (2023). DOI: [10.48550/arxiv.2306.02732](https://doi.org/10.48550/arxiv.2306.02732) . DOI links to the corresponding arXiv preprint, not the proceedings version. [Online source](#)
- Zaffran, M., Féron, O., Goude, Y., Josse, J., Dieuleveut, A.: Adaptive conformal predictions for time series. In: *Proceedings of the 39th International Conference on Machine Learning. Proceedings of Machine Learning Research*, vol. 162, pp. 25834–25866 (2022). DOI: [10.48550/arxiv.2202.07282](https://doi.org/10.48550/arxiv.2202.07282) . DOI links to the corresponding arXiv preprint, not the proceedings version. [Online source](#)
- Zhang, S., Guo, B., Dong, A., He, J., Xu, Z., Chen, S.X.: Cautionary tales on air-quality improvement in beijing. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences* **473**(2205), 20170457 (2017)
DOI: [10.1098/rspa.2017.0457](https://doi.org/10.1098/rspa.2017.0457)

Statements and Declarations

Data and code availability

The source data are available from UCI at DOI: [10.24432/C5RK5G](https://doi.org/10.24432/C5RK5G) under CC BY 4.0. The fixed analysis-code snapshot v0.3.0 is archived at DOI: [10.5281/zenodo.23171630](https://doi.org/10.5281/zenodo.23171630). That software DOI contains code, environment requirements, protocols, method documentation and checksums; it excludes data, fitted models, predictions and the manuscript.

The complete research materials are distributed through the [GitHub research releases](#). The full archive contains source data and partitions in `outputs/02_datos_estudio`; original/extended models and predictions in `outputs/03_experimento` and `outputs/05_revision`; and the manuscript, figures and bibliography in `outputs/04_manuscrito_springer`. The revised package adds the factorial, random-assignment scripts, predictions and diagnostics in `outputs/06_factorial`. This extension reuses the stored models without refitting and is not included in the v0.3.0 DOI. The smaller `source_and_tables.zip` provides source, documentation and tables. SHA-256 manifests identify exact contents. Original results remain traceable to v0.3.0 commit `a8364ec884bbf45bf754053762d935497a68125f`. The software DOI is not an article DOI or a DOI for the complete research package. MIT applies to original code; data and Springer template files retain their separate licenses.

AI use

OpenAI Codex desktop (recorded model identifier: `gpt-6-astra`) assisted during 2–6 October 2026 with literature search and organization, study-design development, code generation and revision, computational execution through Python and R tools, automated verification, figure preparation from numerical outputs, and manuscript drafting and editing. Further analyses and editorial revisions requested on 7 October 2026 used the same tool. Local agent-session records also support the limited chronology audit described above; no immutable backend build identifier was recorded. The author confirmed review and verification of the methods, code, results, figures, references and conclusions, approved the preceding manuscript, and accepts responsibility for the content. The additional analyses reported in this revision require renewed author review before submission. Computational cross-checks are distinguished from independent human reproduction. The AI system is not an author.

Author contributions

Renato Quispe Vargas reviewed the methods, analysis code, results, figures and references and approved the preceding manuscript.

Funding

This research received no funding.

Competing interests

The author declares no competing interests.

Ethics

The analysis uses publicly distributed environmental measurements and contains no participant-level health records. Any institution-specific determination or statement required by the selected journal remains for the authors to confirm.